

Validate Agentic AI Before It Goes Live.

Agents take actions, can you prove their safety?

THE CHALLENGE

1 Specification Gap

Intent is buried in prompts, docs, and tribal knowledge — nothing testable.

79%

of organizations lack written AI-agent policies

[EMA Research](#)

2 Action Gap

Tools, APIs, memory, and RAG create hidden action paths agents can take.

53%

reported agent permission drift

[CSA](#)

3 Evidence Gap

No pre-live proof exists for security sign-off or release approvals.

1 in 3

had an AI-agent security incident


[Verizon CISO Report](#)

65% of those incidents caused real business impact — the cost of finding out in production. [CSA Autonomous Agents Report](#)

THE SOLUTION

NuGuard: the Policy-to-Production Security Pipeline

NuGuard closes all three gaps with an automated pipeline between source code and production - confidently deploy your AI Application in production.



Build AI trust before launch

NuGuard AI takes minutes to deploy, and supports the most secure deployments (air-gapped).

4

AI SECURITY CAPABILITIES

20+

FRAMEWORKS SUPPORTED

1000+

DYNAMIC TESTS

0

LLM KEYS REQUIRED · AIR-GAPPED READY

What NuGuard delivers

01 DISCOVER — AI-SBOM

Map agents, models, tools, prompts, datastores, guardrails, auth/IAM, IaC, and PII/PHI directly from source code.

02 ANALYZE + DEFINE

Detect risks in the application stack and the supply chain, map them into known CVEs, and MITRE framework.

03 VALIDATE + RED TEAM

Run policy-driven static/dynamic checks, MCP/tool validation, prompt injection, tool abuse, and exfiltration tests.

04 GOVERN — EVIDENCE

Generate pass/fail evidence, risk scores, remediation guidance, compliance mappings, and audit ready reports.

Business outcomes

Security posture	Unknown agent behavior becomes documented, testable risk.
Release velocity	CI/CD evidence cuts late-stage security churn.
Incident response	Policy-linked findings speed root cause and remediation.
Regulatory readiness	Artifacts align to NIST AI RMF, EU AI Act, MITRE ATLAS, OWASP LLM Top 10.
Enterprise trust	Evidence-backed claims for boards, customers, procurement, and M&A.
No lock-in	Open-source CLI, Python API, GitHub Actions; most checks need no LLM key.

Why NuGuard is different

Pre-production first

Behavior is validated before production traffic — not patched after an incident.

Redteam Innovation Engine

1000+ tests generated on-the-fly for your specific AI stack and Behavior Policy

Audit-Ready Evidence Package

Every run produces AI-SBOM, supply-chain analysis, compliance mapping, and sandbox validation