

Validate Agentic AI Before It Goes Live.

Agents take actions, can you prove their safety?

THE CHALLENGE

1 Specification Gap

Intent is buried in prompts, docs, and tribal knowledge — nothing testable.

79%

of organizations lack written AI-agent policies
[EMA Research](#)

2 Action Gap

Tools, APIs, memory, and RAG create hidden action paths agents can take.

53%

reported agent permission drift
[CSA](#)

3 Evidence Gap

No pre-live proof exists for security sign-off or release approvals.

1 in 3

had an AI-agent security incident
[Varion CISO Report](#)

65%

of those incidents caused real business impact — the cost of finding out in production. [CSA Autonomous Agents Report](#)

THE SOLUTION

NuGuard: the Policy-to-Production Security Pipeline

NuGuard closes all three gaps with an automated pipeline between source code and production - confidently deploy your AI Application in production.

```
graph LR; SC[Source code] --> D[01 · DISCOVER  
AI stack discovery  
(AI_SBOM)]; D --> F[02 · DEFINE  
Behavior policy]; F --> V[03 · VALIDATE  
Automated red teaming  
in sandbox (inc.  
supply-chain analysis)]; V --> G[04 · GOVERN  
Audit evidence &  
release authorization]; G --> P([PRODUCTION  
evidence-backed]); V -.-> R[Remediate & re-validate  
implement recommendations]; R -.-> SC;
```

The diagram illustrates the NuGuard Policy-to-Production Security Pipeline. It begins with 'Source code' (represented by a code icon) which flows into a sequence of four stages: 01 · DISCOVER (AI stack discovery (AI_SBOM)), 02 · DEFINE (Behavior policy), 03 · VALIDATE (Automated red teaming in sandbox (inc. supply-chain analysis)), and 04 · GOVERN (Audit evidence & release authorization). The final stage leads to 'PRODUCTION evidence-backed'. A feedback loop is shown where the '03 · VALIDATE' stage triggers a 'Remediate & re-validate' step (implement recommendations), which then feeds back into the 'Source code'.

SHIFT AGENT RISK LEFT — validate before live traffic.

Continued: capabilities & outcomes →

Build AI trust before launch

NuGuard AI takes minutes to deploy, and supports the most secure deployments (air-gapped).

4

AI SECURITY CAPABILITIES

20+

FRAMEWORKS SUPPORTED

1000+

DYNAMIC TESTS

0

LLM KEYS REQUIRED · AIR-GAPPED READY

What NuGuard delivers

01 DISCOVER — AI-SBOM

Map agents, models, tools, prompts, datastores, guardrails, auth/IAM, IaC, and PII/PHI directly from source code.

02 ANALYZE + DEFINE

Detect risks in the application stack and the supply chain, map them into known CVEs, and MITRE framework.

03 VALIDATE + RED TEAM

Run policy-driven static/dynamic checks, MCP/tool validation, prompt injection, tool abuse, and exfiltration tests.

04 GOVERN — EVIDENCE

Generate pass/fail evidence, risk scores, remediation guidance, compliance mappings, and audit ready reports.

Business outcomes

Security posture	Unknown agent behavior becomes documented, testable risk.
Release velocity	CI/CD evidence cuts late-stage security churn.
Incident response	Policy-linked findings speed root cause and remediation.
Regulatory readiness	Artifacts align to NIST AI RMF, EU AI Act, MITRE ATLAS, OWASP LLM Top 10.
Enterprise trust	Evidence-backed claims for boards, customers, procurement, and M&A.
No lock-in	Open-source CLI, Python API, GitHub Actions; most checks need no LLM key.

Why NuGuard is different

Pre-production first

Behavior is validated before production traffic — not patched after an incident.

Redteam Innovation Engine

1000+ tests generated on-the-fly for your specific AI stack and Behavior Policy

Audit-Ready Evidence Package

Every run produces AI-SBOM, supply-chain analysis, compliance mapping, and sandbox validation