

BEHAVIOR VALIDATION FOR AGENTIC AI

Validate Agentic AI Before It Goes Live.

Agents take actions, can you prove their safety?
NuGuard finds, tests, and governs that behavior before it reaches production.

THE CHALLENGE

1 Specification Gap

Intent is buried in prompts, docs, and tribal knowledge — nothing testable.

79%

of organizations lack written AI-agent policies

EMA Research

2 Action Gap

Tools, APIs, memory, and RAG create hidden action paths agents can take.

53%

reported agent permission drift

CSA

3 Evidence Gap

No pre-live proof exists for security sign-off or release approvals.

1 in 3

had an AI-agent security incident

Verizon CISO Report

65%


of those incidents caused real business impact — the cost of finding out in production.

CSA Autonomous Agents Report

THE SOLUTION

NuGuard: the Policy-to-Production Security Pipeline

NuGuard closes all three gaps with an automated pipeline between source code and production - confidently deploy your AI Application in production.



Build AI trust before launch

NuGuard AI is a pre-production behavioral validation platform for agentic AI: it discovers agent architecture, enforces governance policy, and generates audit-ready evidence before deployment.

4

AI SECURITY CAPABILITIES

20+

FRAMEWORKS SUPPORTED

1000+

DYNAMIC TESTS

0

LLM KEYS REQUIRED · AIR-GAPPED READY

What NuGuard delivers

01 DISCOVER — AI-SBOM

Map agents, models, tools, prompts, datastores, guardrails, auth/IAM, IaC, and PII/PHI directly from source code.

02 ANALYZE + DEFINE

Detect risks in the application stack and the supply chain, map them into known CVEs, and MITRE framework.

03 VALIDATE + RED TEAM

Run policy-driven static/dynamic checks, MCP/tool validation, prompt injection, tool abuse, and exfiltration tests.

04 GOVERN — EVIDENCE

Generate pass/fail evidence, risk scores, remediation guidance, compliance mappings, and audit ready reports.

Business outcomes

Security posture	Unknown agent behavior becomes documented, testable risk.
Release velocity	CI/CD evidence cuts late-stage security churn.
Incident response	Policy-linked findings speed root cause and remediation.
Regulatory readiness	Artifacts align to NIST AI RMF, EU AI Act, MITRE ATLAS, OWASP LLM Top 10.
Enterprise trust	Evidence-backed claims for boards, customers, procurement, and M&A.
No lock-in	Open-source CLI, Python API, GitHub Actions; most checks need no LLM key.

Why NuGuard is different

Pre-production first

Behavior is validated before production traffic — not patched after an incident.

Redteam Innovation Engine

1000+ tests generated on-the-fly for your specific AI stack and Behavior Policy

Audit-Ready Evidence Package

Every run produces AI-SBOM, supply-chain analysis, compliance mapping, and sandbox validation